
Bouncer: Runtime Competence Auditing for Learned Systems Controllers

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Learned systems controllers can outperform fixed heuristics on their training distri-
2 bution yet silently regress after workload shift. Input-distribution monitors cannot
3 determine whether a controller’s decisions still improve the system. We present
4 BOUNCER, a runtime trust gate that measures *realized competence*: a secret sample
5 of resource partitions runs the learned controller, another runs a vetted fallback,
6 and a sequential detector gates the remaining partitions when the learned-minus-
7 fallback reward falls below a threshold. This requires representative sampling,
8 set-local reward, and policy contrast that survives assignment. In a new trained-
9 controller case study, an offline logistic cache-insertion policy validates at 96.3%
10 and initially beats LRU, but a label-semantic shift reverses its advantage while
11 leaving its input histogram exactly unchanged. Across 12 seeds, BOUNCER detects
12 all reversals in one audit window, retains 87.5% of the clean learned gain, and
13 recovers 87.5% of the loss toward LRU. The study is controlled trace replay rather
14 than full-system validation; separate ChampSim experiments expose where reward
15 locality and stateful reassignment fail. The artifact is deterministic and released.

16 1 Why monitor competence?

17 Learned cache replacement, prefetching, memory scheduling, and branch prediction can beat hand-
18 designed policies by adapting to workload structure [2, 5, 6, 10]. Their deployed advantage is
19 nevertheless conditional: a model trained on one program or phase may become worse than the
20 heuristic it displaced. Aggregate counters show that performance changed, but not whether the
21 learned policy caused the change. Feature-based out-of-distribution (OOD) detection is also indirect:
22 a concept shift can preserve inputs while reversing which action is useful [3, 4].

23 Suppose a learned controller C has a vetted fallback π_0 . For audit window w , define the competence
24 advantage

$$\Delta_w = \bar{r}_w^C - \bar{r}_w^{\pi_0}, \quad r \in [0, r_{\max}], \quad (1)$$

25 where reward is the system outcome attributable to a decision, such as a cache hit. We want to deploy
26 C while $\Delta_w \geq \tau$ and otherwise fall back. The counterfactual is unavailable: one physical decision
27 cannot simultaneously run both policies.

28 2 Secret online A/B auditing

29 **Estimator and gate.** BOUNCER repurposes hardware set dueling [8]. It secretly assigns n_C cache
30 sets to C , n_F to π_0 , and the rest to a follower pool (Figure 1). Unlike a global before/after comparison,
31 both leaders observe the same concurrent workload. Their measured difference

$$\hat{\Delta}_w = \bar{r}_{C,w} - \bar{r}_{F,w} \quad (2)$$




Figure 1: BOUNCER measures rather than predicts trust. Secret Leader-C partitions run the learned controller, Leader-F partitions run the fallback, and followers obey the gate. The mechanism is faithful only when the sampled reward represents deployed traffic and remains attributable to the sampled partition.

32 feeds a lower CUSUM $S_w = \max\{0, S_{w-1} + K - \hat{\Delta}_w\}$, which gates followers when $S_w > H$ [7].
 33 Hysteresis and measured probes support re-trust. Leaders remain policy-fixed so the gate continues
 34 to observe both arms.

35 **When does the comparison mean what it says?** Three conditions are load-bearing. **Representa-**
 36 **tative sampling** requires the leader estimate to target decision-weighted follower traffic. **Set-locality**
 37 requires a decision and credited reward to share the sampled partition; replacement hits satisfy this,
 38 whereas a prefetch initiated in one set can benefit another. **Assignment-identifiability** requires the
 39 policy contrast to survive leader assignment. A stateless controller satisfies this immediately. A
 40 stateful controller instead needs persistent leaders, shadow-updated state, or a warm-up quarantine.
 41 Secret reshuffling that repeatedly cold-starts one arm can erase the very contrast being measured.

42 **Expected floor.** The full paper proves the following compact accounting result. Over T windows,
 43 partition regret into (i) fully open harmful windows, (ii) the learned-policy exposure retained for
 44 auditing after gating, and (iii) false-alarm switching. If the expected number of fully open windows
 45 per harmful episode is at most D , each set is exposed after gating with marginal probability at most
 46 ϕ , the false-alarm probability per window is at most α , and each false-alarm episode costs at most
 47 c_{sw} , then

$$\mathbb{E}[R_T] \leq N_{ep} D r_{\max} + \phi T_{\text{harm}} r_{\max} + \alpha T c_{sw}. \quad (3)$$

48 The result is expectation-level and conditional on the stated assumptions; it does not make a pointwise
 49 safety claim. Appendix B gives the proof skeleton and separates this accounting guarantee from the
 50 experiment.

51 3 One trained controller under concept shift

52 We add the smallest experiment missing from the original artifact: a controller whose deployment
 53 decisions come from fitted parameters. A logistic model sees one of four program-counter (PC)
 54 classes and predicts whether a line will be reused within 64 future references. Training and validation
 55 labels are computed from independent seeded traces containing recurring and one-shot lines. The
 56 controller inserts predicted-reusable lines and bypasses predicted-streaming lines; ordinary true LRU
 57 is the fallback. The model uses 20,000 training and 5,000 validation examples and achieves 96.28%
 58 validation accuracy.

59 The deployment cache has 64 sets, eight ways, and 64 accesses per set-window. Across 60 windows,
 60 half the accesses target four recurring lines and half are one-shot streams. At window 30, the




Figure 2: Mean and 95% normal confidence intervals over 12 seeds. The trained insertion controller beats LRU before the semantic shift and collapses after it, although the PC histogram is unchanged. The measured leader gap reverses and gates followers one window later.

relationship between PC and future reuse reverses. Crucially, every PC still occurs exactly 25% of the time in every set-window, so a PC-histogram OOD detector has total-variation distance zero and never alarms. We use eight fixed secret leaders per arm and 12 independently seeded deployments. Fixed leaders avoid a cache-state reset confound; this benign-shift study therefore does not test assignment secrecy against an adaptive attacker.

Figure 2 shows the causal reversal. Before the shift, hit rates are 0.4375 for the learned policy, 0.3276 for LRU, and 0.4237 for BOUNCER; hence BOUNCER retains 87.5% of the available learned gain. After the shift, they are 0, 0.3279, and 0.2869. All 12 runs gate one audit window after the shift and recover 87.5% of the learned-to-LRU loss; the remaining gap is exactly the fixed Leader-C exposure. These values are steady summaries that exclude five cold-start/transition windows per regime. They demonstrate the intended mechanism under a controlled concept shift, not generalization to production cache policies or an IPC benefit.

4 Broader evidence and boundaries

The trained case complements, rather than replaces, the artifact’s broader tests. In the seeded competence harness, $\hat{\Delta}$ tracks known ground truth ($R^2 = 0.997$); the released operating point gates in two windows, limits the worst attacked steady window to 0.64% below the fallback, and detects 50/50 input-mimicry attacks while an input-OOD monitor detects 0/50. These are synthetic mechanism tests with explicit reward models.

The ChampSim integration supplies a different kind of evidence: it runs the auditor inside a cycle-level simulator on SPEC traces, but exposes two important boundaries rather than a clean learned-controller win. Prefetch reward is not set-local, and repeatedly reassigned replacement sets carry path-dependent cache state. Fixed replacement leaders recover a measurable gap, while reseeding can attenuate it. Thus the new trained trace replay establishes that fitted decisions can be gated; ChampSim establishes that the abstraction cannot be transferred blindly. Neither experiment measures RTL area, energy, timing, or real silicon. A 16-bit state-count proposal is 406 bytes, but it is not a synthesis result.

Table 1: Each evaluation answers a different question; none substitutes for the others.

Evidence	Strength	Deliberate boundary
Trained trace replay	Fitted controller; 12 seeded concept shifts	Controlled workload, hit rate rather than IPC
Synthetic harness	Known competence, adaptive attacks, exact claim checks	Parametric reward model
ChampSim/SPEC	Cycle-level execution and real cache state	Boundary study; committed logs, no clean learned-controller validation

86 **Relationship to prior work.** OOD and concept-drift detectors ask whether deployment resembles
87 reference data [3, 4]; BOUNCER instead asks whether the controller still beats a fallback. Safe-control
88 architectures such as Simplex and learned-policy shields switch to a trusted controller when a model
89 or safety condition fails [1, 9]. Our narrower systems setting replaces a verified plant invariant
90 with a randomized online performance comparison. Set dueling supplies the comparison substrate,
91 while sequential detection supplies the decision rule. The contribution is their use as a competence
92 certificate with explicit locality and state conditions, not a new cache-replacement learner.

93 **5 Limitations and conclusion**

94 The controlled controller is intentionally small, its PC/reuse reversal is constructed, and hit rate is not
95 IPC. Fixed leaders preserve state but provide less freshness than epoch reshuffling, and the trained
96 experiment evaluates benign shift rather than a leader-aware adversary. More broadly, BOUNCER
97 cannot faithfully audit nonlocal rewards, unrepresentative traffic, sub-resolution harm, or controllers
98 whose relative behavior is destroyed by assignment. These are scope boundaries, not details hidden
99 by the detector.

100 Within that scope, the result is encouraging: online randomized comparison can retain most of a
101 trained controller’s upside while recovering most of a known fallback when competence reverses,
102 even when a simple input monitor has no signal. The released experiment and claim checker make
103 that workshop-sized result directly reproducible, while the longer artifact retains the proofs, adaptive
104 attacks, and full sensitivity studies.

References

- [1] Mohammed Alshiekh, Roderick Bloem, Rüdiger Ehlers, Bettina Könighofer, Scott Niekum, and Ufuk Topcu. Safe reinforcement learning via shielding. In *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence (AAAI)*, pages 2669–2678. AAAI Press, 2018.
- [2] Rahul Bera, Konstantinos Kanellopoulos, Anant V. Nori, Taha Shahroodi, Sreenivas Subramoney, and Onur Mutlu. Pythia: A customizable hardware prefetching framework using online reinforcement learning. In *Proceedings of the 54th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, pages 1121–1137, 2021. doi: 10.1145/3466752.3480114.
- [3] João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):44:1–44:37, 2014. doi: 10.1145/2523813.
- [4] Dan Hendrycks and Kevin Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- [5] Akanksha Jain and Calvin Lin. Back to the future: Leveraging belady’s algorithm for improved cache replacement. In *Proceedings of the 43rd International Symposium on Computer Architecture (ISCA)*, pages 78–89, 2016. doi: 10.1109/ISCA.2016.17.
- [6] Evan Zheran Liu, Milad Hashemi, Kevin Swersky, Parthasarathy Ranganathan, and Junwhan Ahn. An imitation learning approach for cache replacement. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, pages 6237–6247, 2020.
- [7] E. S. Page. Continuous inspection schemes. *Biometrika*, 41(1/2):100–115, 1954. doi: 10.1093/biomet/41.1-2.100.
- [8] Moinuddin K. Qureshi, Aamer Jaleel, Yale N. Patt, Simon C. Steely, and Joel Emer. Adaptive insertion policies for high performance caching. In *Proceedings of the 34th Annual International Symposium on Computer Architecture (ISCA)*, pages 381–391. ACM, 2007. doi: 10.1145/1250662.1250709.
- [9] Lui Sha. Using simplicity to control complexity. *IEEE Software*, 18(4):20–28, 2001. doi: 10.1109/MS.2001.936213.
- [10] Zhan Shi, Xiangru Huang, Akanksha Jain, and Calvin Lin. Applying deep learning to the cache replacement problem. In *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, pages 413–425, 2019. doi: 10.1145/3352460.3358319.

136 A Trained-controller protocol

137 The released script `experiments/exp_trained_controller.py` owns the entire trained experi-
138 ment. The model is a regularized binary logistic regressor with an intercept and four one-hot PC
139 features. Full-batch gradient descent runs for 800 steps at learning rate 0.4 with 10^{-3} weight decay.
140 Training traces contain equal numbers of recurring and one-shot references; 3% of examples cross
141 the usual PC/reuse association. A label is one exactly when the same tag appears within the next 64
142 references. Training and validation traces use different seeds.

143 Deployment uses a set-associative trace replay with true LRU ordering. On a miss, the learned policy
144 inserts the line if its fitted reuse probability is at least 0.5 and otherwise bypasses it. The fallback
145 always inserts. Hot and stream references are balanced, as are all four PC classes. The PC classes
146 associated with hot and stream references exchange at window 30. Hot tags rotate each window,
147 preventing the learned controller from coasting indefinitely on state established before the shift.

148 Each evaluation seed draws one fixed secret partition: eight learned leaders, eight fallback lead-
149 ers, and 48 followers. The one-sided lower CUSUM uses $K = 0.02$ and $H = 0.12$. Evidence
150 from window w changes follower routing at window $w + 1$, so the reported one-window delay
151 is causal rather than same-window routing. The experiment asserts detection in every seed, an
152 unchanged PC histogram, a positive pre-shift learned advantage, a negative post-shift advantage,
153 and minimum retained/recovered utility. The JSON and all manuscript numbers are checked by
154 `scripts/check_ml4sys_numbers.py`.

155 B Expected-floor proof skeleton

156 Let $R_T = \sum_{w=1}^T (\bar{r}_w^{\pi_0} - \bar{r}_w^{\text{BOUNCER}})$. Split windows into three disjoint charges. During fully open
157 harmful windows, bounded reward charges at most $r_{\max}$ per window; the assumed expected count
158 is $N_{\text{ep}}D$. After gating, each harmful decision remains on C only through the audit exposure.
159 Conditional on traffic and history, secret uniform sampling gives every set marginal exposure at
160 most ϕ , so linearity of expectation charges at most $\phi T_{\text{harm}} r_{\max}$ without requiring independent
161 traffic. Finally, the expected number of false alarms is at most αT , and each episode costs at most
162 c_{sw} . Summing the three charges yields Equation 3; windows where C outperforms π_0 contribute
163 nonpositive regret and may be discarded. The premise D includes both initial detection and any false
164 re-trust. Only the exposure term has a separate high-probability refinement in the full artifact.

165 C Reproduction and claim ownership

166 Run `./run_ml4sys.sh`. It regenerates the training/evaluation JSON and figure, checks 13 headline
167 values, compiles this PDF, and verifies that the main text ends by page four and references begin on
168 page five. All random generators are explicitly seeded. The trained experiment requires NumPy and
169 Matplotlib; the repository pins the complete environment. The full synthetic and ChampSim artifact
170 remains available through `./run_all.sh`; the latter rebuilds figures from committed simulator logs
171 and does not rerun every SPEC simulation.